

Threading Building Blocks

Martin Kruliš





Introduction

- Now part of Intel oneAPI
 - <https://spec.oneapi.io/versions/latest/elements/oneTBB/source/nested-index.html>
- J. Reinders - Intel® Threading Building Blocks: Outfitting C++ for Multi-core Processor Parallelism – O'REILLY 2007
 - Somewhat obsolete





Introduction

- C++ library
 - Namespace `oneapi::tbb`
- Templates
- Compatible with other threading libraries (pthreads, OpenMP,...)
- Works with tasks, not threads

Used to be `tbb` only



Contents of the Library

Parallel
Algorithms

Flow Graph

Containers

Task Scheduler

Thread Local
Storage

Threads

Sync.
Primitives

Utilities
(Timing,...)

Memory
Allocator



Initialization

- None required (starts automatically)
 - `task_scheduler_init` class
 - Constructor takes number of threads
 - DEPRECATED!
 - `global_control` class
 - Constructor takes limit values for number of threads and their stack sizes

Optional, for debugging/testing only



Parallel algorithms

- `parallel_for`
- `parallel_for_each`
- `parallel_reduce`, `parallel_deterministic_reduce`
- `parallel_scan`
- `parallel_do`
- `pipeline`, `parallel_pipeline`
- `parallel_sort`
- `parallel_invoke`



parallel_for (type 1)

```
template<typename Range, typename Body>  
void parallel_for( const Range& range,  
    const Body& body );
```

compare to

```
template <class InIter, class UnaryFunc>  
UnaryFunc for_each(InIter first, InIter  
    last, UnaryFunc f);
```

Range vs. Iterator



Splittable Concept

- A splittable object has the following constructor:
`X::X(X& x, Split)`
- Unlike copy constructor, the first argument is not constant
- Divides (splits) the first argument into two parts
 - one is stored back into the first argument
 - other is stored in the newly constructed object
- Applies to both Range and Body
 - splitting of a range into two parts (first part into argument, second part into newly created object)
 - splitting body to two instances executable in parallel



Body

- Action (functor) to be executed in parallel

```
void Body::operator() (Range& range) const
```

- Unlike STL functor, the argument is a range
 - your job is to iterate the range



Range

- `bool R::empty() const`
- `bool R::is_divisible() const`
- `R::R(R& r, split)`
- split should create two parts of similar size
 - recommendation, not requirement

```
struct TrivialIntegerRange {  
    int lower, upper;  
    bool empty() const { return lower == upper; }  
    bool is_divisible() const { return upper > lower+1; }  
    TrivialIntegerRange(TrivialIntegerRange& r, split) {  
        int m = (r.lower+r.upper) / 2;  
        lower = m; upper = r.upper; r.upper = m;  
    }  
};
```



blocked_range<Value>

```
template<typename Value> class blocked_range;
```

- Value needs to support:
 - copy constructor, destructor
 - < (Value, Value)
 - + (Value i, size_t k) – k-th value after i
 - - (Value, Value) – returns size_t – distance
- Grainsize
 - minimal number of elements in a range
- Function `begin()`, `end()`
 - minimal and maximal value – half-closed interval



parallel_for (type 1) – Again

- Splits the range until it cannot be split further
 - `Range::is_divisible()`
- Creates a copy of body for each subrange
- Executes bodies over ranges
- Actual splitting and execution is done in parallel in a more sophisticated manner
 - Better use of CPU caches



blocked_range - Grainsize

- How to determine proper grainsize
 - depends on the actual algorithm
 - one call ~ at least 10.000 to 100.000 instructions
 1. set the grainsize to ~10.000
 2. run on one processor
 3. halve the grainsize
- Repeat steps 2 and 3 and observe slowdown
 - the performance decreases with growing overhead (1 core)
 - optimal grainsize is considered for 5-10% slowdown
- No exact way to get the best value





Partitioners

- Other solution to grainsize problem
- Third (optional) parameter to `parallel_for`
- Range may not be split to the smallest parts
 - `simple_partitioner` (default) – split to the smallest parts
 - `auto_partitioner` – enough splits for load balancing, may not provide optimal results
 - `affinity_partitioner` – same as `auto_partitioner` but better cache affinity



parallel_for (type 2)

```
template<typename Index, typename Func>
void parallel_for(Index first, Index last,
    Index step, const Func& f);
```

```
template<typename Index, typename Func>
void parallel_for(Index first, Index last,
    const Func& f);
```

- Index must be an integral type

- Semantics:

```
for (auto i=first; i<last; i+=step) f(i);
```



parallel_reduce<Range,Body>

```
template<typename Range, typename Body>  
void parallel_reduce(const Range& range,  
    const Body& body) ;
```

- Similar to parallel for, but returns a result
- Example: sum of all values
- New requirement for Body

```
void Body::join(Body& rhs) ;
```

- join two results into one



parallel_reduce<Range,Body>

```
template<typename Range, typename Body>
void parallel_reduce(const Range& range,
    const Body& body);
```

- One body may process more ranges (seq.)
- `operator()` and `join()` can run parallel with splitting constructor (but not each other)
 - does not cause problems in most cases
- Noncommutative operation \oplus
 - `left.join(right)` should produce `left \oplus right`



parallel_reduce (type 2)

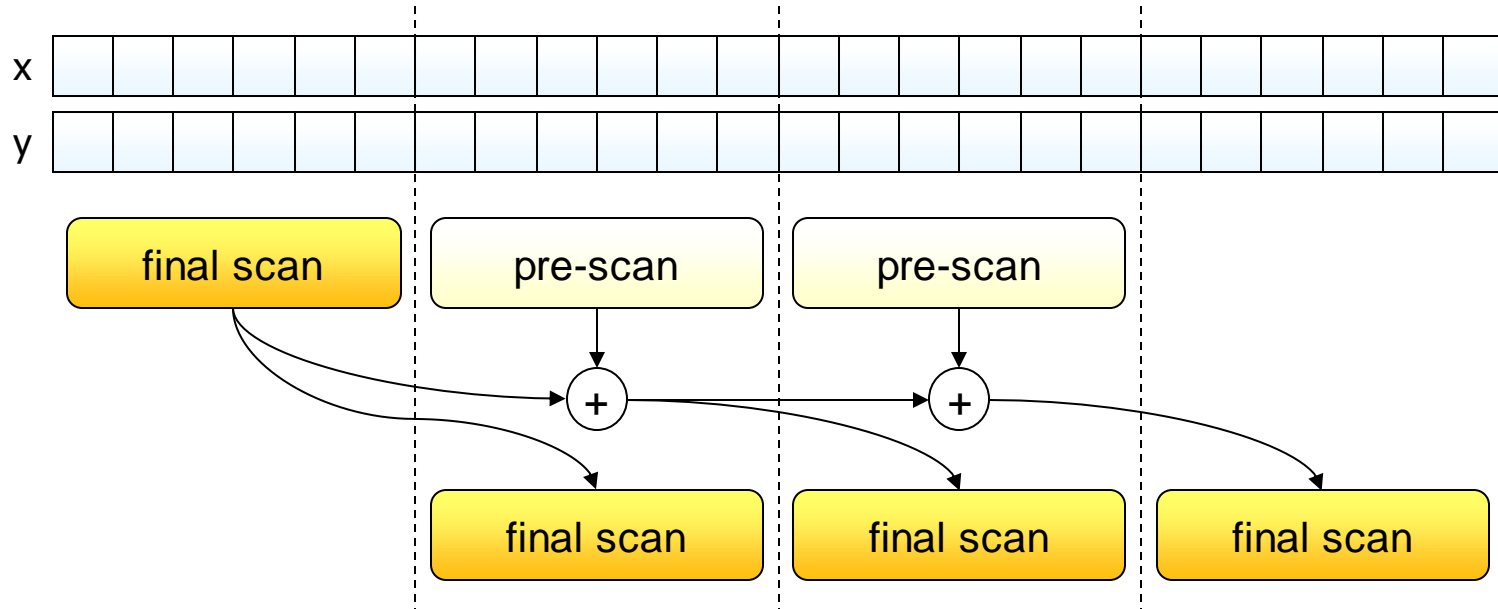
```
template<typename Range, typename Value,  
        typename Func, typename Reduction>  
void parallel_reduce(const Range& range,  
                    const Value& identity, const Func& f,  
                    const Reduction& reduction);
```

- Identity is the left identity element for left operator of `Func::operator()`
- Value `Func::operator()` (`const Range& range`, `const Value& x`) accumulates subrange to the result starting with value `x`
- Value `Reduction::operator()` (`const Value& x`, `const Value& y`)
- Lambda-friendly, more data copying



parallel_scan<Range,Body>

- $y_0 = \text{id}_{\oplus} \oplus x_0$, $y_i = y_{i-1} \oplus x_i$
- \oplus needs to be associative





parallel_for_each

- Like STL's `for_each`
 - Same semantics, except for parallel processing
- Different Body than e.g. `parallel_for`
 - `operator()` is unary and must be reentrant
 - Optional 2nd argument is feeder, which is used to add more work that is process by subsequent calls (for better scalability)
 - no splitting
- Not suitable for trivial operations
 - less than ~ 10.000 instructions in `Body::operator()`



pipeline

- Sequence of filters

```
class pipeline {  
public:  
    void add_filter(filter& f);  
    void run(size_t max_num_of_live_tokens);  
};
```

- Filters are serial or parallel

- serial is either in-order or out-of-order

- Filters derive from class filter and override virtual

```
void* operator()(void *item);
```

- the first filter's operator() is called repeatedly until it returns NULL
- the results returned by the last filter are ignored



parallel_pipeline

- Strongly typed lambda-friendly pipeline

```
void parallel_pipeline(  
    size_t max_number_of_live_tokens,  
    const filter_t<void, void>& filter_chain);
```

- Filters

- combined using & operator
- same modes as original pipeline
- input and output type (first and last are voids)
- functor for the action



parallel_pipeline

```
float RootMeanSquare(float* first, float* last) {
    float sum=0;
    parallel_pipeline(/*max_number_of_live_token=*/16,
        make_filter<void, float*>(filter::serial,
            [&](flow_control& fc) -> float* {
                if (first < last) return first++;
                else { fc.stop(); return nullptr; }
            }
        ) &
        make_filter<float*, float>(filter::parallel,
            [](float* p) { return (*p)*(*p); }
        ) &
        make_filter<float, void>(filter::serial,
            [&](float x) { sum += x; }
        )
    );
    return sqrt(sum);
}
```



Flow graph

- Generalization of the pipeline
 - arbitrary oriented graph
 - nodes and explicit (dataflow) connections
- More complex rules for execution
 - e.g., no explicit “number of live tokens”
- Pre-defined set of node types
 - function, join, limiter,...
- See the documentation...



parallel_sort

```
template<typename RndAccIter>
void parallel_sort(RndAccIter begin, RndAccIter end);
template<typename RndAccIter, typename Compare>
void parallel_sort(RndAccIter begin, RndAccIter end,
    const Compare& comp);
```

- Parallel unstable sort
- Average time complexity of $O(N \log (N))$
- Usage:

```
parallel_sort(a, a + N);
parallel_sort(b, b+N, std::greater<float>());
```



parallel_invoke

- Parallel invocation of multiple functors
 - May be used to invoke completely different tasks in a simple manner

```
tbb::parallel_invoke(  
    myfunc1,  
    []{ bar(2); },  
    []{ bar(3); },  
);
```



Containers

```
concurrent_unordered_map<Key, Val, Hash, Eq, Alloc>  
concurrent_unordered_set<Key, Hash, Eq, Alloc>  
concurrent_hash_map<Key, T, HashCompare, Alloc>  
concurrent_queue<T, Alloc>  
concurrent_bounded_queue<T, Alloc>  
concurrent_priority_queue<T, Compare, Alloc>  
concurrent_vector<T>
```

- Selected operations are thread-safe
- Can be used with “raw” threads, OpenMP,...
- Designed for high concurrency
 - fine grained locking, lock-free algorithms
- Higher overhead than STL (due to concurrency)



concurrent_vector

- Concurrent growth and access
- Careful with exceptions
 - Special rules for item constructors and destructors
- Existing elements are never moved
 - iterators and references not invalidated by growth
 - results in fragmentation – call `compact()`
- Copy construction and assignment are not thread-safe



concurrent_vector

- `grow_by(size_type delta [, const T& t]);`
 - adds `delta` elements (atomically) to the end of vector
 - returns original size of the vector ... `k`
 - new elements have indexes `[k, k+delta)`
- `size_t push_back(const T& value);`
 - atomically adds copy of value to the end of the vector
 - returns index of the new value
- `access`
 - `operator[]`, `at`, `front`, `back` – all const and non-const
 - can be called concurrently while the vector grows
 - may return reference that is currently being constructed!



concurrent_vector

- range (const and non-const)
 - used in parallel algorithms
- size()
 - number of elements, including those currently being constructed
- standard ISO C++ random access iterators



concurrent_queue

- push, try_pop
 - non-blocking
- No STL-like front and back
 - can't be made thread-safe
 - unsafe_begin, unsafe_end, unsafe_size
- concurrent_bounded_queue
 - added later, originally concurrent_queue was bounded
 - limits the maximal number of elements in queue
 - push and pop are blocking
 - active lock used for blocking operations – should not wait long!
 - size() returns number of pushes minus number of pops



concurrent_queue

- Order of items partially preserved
 - if one threads pushes X before Y and another thread pops both values, then X is popped before Y
- Provides iterators
 - for debugging only, invalidated on modification
 - STL compliant iterator and const_iterator
- To be used wisely
 - maybe parallel_do or pipeline can be used as well
 - parallel algorithms perform better – avoid one bottleneck (the queue), better use caches, ...
 - resist the temptation to implement producer-consumer (or similar) algorithms



concurrent_priority_queue

- Similar to STL priority_queue
 - With fixed underlying container
 - try_pop() instead of pop()
 - thread-unsafe empty() and size()



concurrent_unordered_map

- Similar to `std::unordered_map`
- Insertion and iteration are safe
 - Even combined
 - Insertion does not invalidate iterators
- Erasing is unsafe
- No visible locking
- Same analogy goes for `unordered_set`



concurrent_hash_map

```
template<typename Key, typename T,  
        typename HashCompare, typename Alloc>  
class concurrent_hash_map;
```

- Provides accessors
 - `const_accessor` and `accessor`
 - smart pointer with implicit lock (read or RW lock)
 - cannot be assigned or copy constructed
 - operators `*` and `->`
 - explicit release of locks – call `release()`
 - if points “to nothing” – can be tested by `empty()`



concurrent_hash_map

- Concurrent operations
 - `find` (returns const or non-const accessor)
 - `insert` (returns const or non-const accessor)
 - `erase`
 - copy constructor
 - the copied table may have operations running on it
- `range(size_t grainsize)`
 - returns (constant or non-constant) range that can be used for parallel algorithms to iterate over contents of the map
 - cannot be run concurrently with other operations
- Forward iterators – `begin()`, `end()`



Memory Allocation

- Optional part of the library
 - C++ and STL compliant allocators
 - C functions (like malloc and free)
 - Memory resources
 - Based on C++17 `std::pmr::polymorphic_allocator`
 - Implementation of `std::pmr::memory_resource`
 - `cache_aligned_resource`
 - `scalable_memory_resource`



Memory Allocation

- `scalable_allocator`
 - each thread has its own memory pool
 - no global lock, prevents false sharing of private memory
- `cache_aligned_allocator`
 - all functionality of `scalable_allocator`
 - aligns memory to cache lines
 - prevents false sharing of separate allocations
 - always pads data to cache line size (usually 128bytes) !
 - use only when it proves to be needed
- `tbb_allocator`
 - uses `scalable_allocator` if possible (TBB malloc is present) or the standard malloc/free



Mutexes

- Many variations
 - different implementation (OS / user space, ...)
 - different semantics (simple/recursive, fair/unfair, plain/read-write)
- Scoped locking pattern
 - acquired lock is represented by an object
 - no explicit unlocking – less prone to errors
 - exception safe



Thread Local Storage

- Separate variable for each thread
 - lazily created (on demand)
- `combinable<T>`
 - reduction variable, each thread has a private copy
 - combined using a bin. functor or unary (`for_each`)
- `enumerable_thread_specific<T>`
 - container with element for each thread (lazy)
 - can access either local value or iterate all values
 - may be combined using a binary functor



Timing

- Thread-safe timing routines
- Classes `tick_count` and `tick_count::interval_t`

```
tick_count t0 = tick_count::now();  
... action being timed ...  
tick_count t1 = tick_count::now();  
printf("time for action = %g seconds\n",  
      (t1-t0).seconds());
```



Task Scheduler

- Low level task engine
 - Manages thread pool, maps tasks on threads
 - Tasks are not preemptive (should not block)
 - No guarantee that potentially parallel tasks will actually run concurrently!
 - Other algorithms are implemented using tasks
 - Can be used directly
 - A bit tedious, algorithms are better



Task Scheduler

- **task_group**
 - Group of tasks being concurrently executed
 - Provides interface to spawn more tasks
 - **run(Func&& f)**
 - **run_and_wait(const Func& f)**
 - **task_handle defer(Func&& f)**
 - Create but not execute a task
 - The handle can be passed to run or run_and_wait
 - **wait()**
 - **cancel()**



Task Scheduler

- `task_arena`
 - Constructor/initializer can specify level of concurrency and priority
 - `enqueue(Func&& f)`
 - Enqueue and return immediately
 - Also can enqueue existing tasks (by handle)
 - `auto execute(Func&& f) -> decltype(f())`
 - Execute and wait for completion
 - Current thread may join arena (optimization) to start immediate execution of the task



Task Scheduler

- Resumable tasks
 - Allows to suspend and later resume task without blocking a thread (e.g., for I/O)

```
oneapi::tbb::task::suspend(  
    [&] (oneapi::tbb::task::suspend_point tag) {  
        usr_async_activity.submit(tag);  
    }  
);  
...
```

I/O, network, GPU task ...

```
oneapi::tbb::task::resume(tag);
```


Once the user-managed
async task concludes




Task Scheduling

- Each thread has its own pool of ready tasks
 - task is ready when it has no children
 - new tasks are spawned to the front of the stack
 - tasks are stolen from the end of the stack


Worker #1



Worker #2



Worker #3



Worker #4

